

Classify the email using the binary classification method. Email Spam detection has two states: a) Normal State – Not Spam, b) Abnormal State – Spam. Use K-Nearest Neighbors for classification. Analyze their performance. Dataset link: The emails.csv dataset on the Kaggle <https://www.kaggle.com/datasets/balaka18/email-spam-classification-dataset-csv>

```
In [16]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, accuracy_score
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [17]: # Load the dataset
df = pd.read_csv("C:/Users/Atharva/OneDrive/Desktop/LP3 code/emails.csv")
print(df.head())
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	\
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	0	
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	0	
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	0	
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	0	
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	0	

	valued	lay	infrastructure	military	allowing	ff	dry	Prediction
0	0	0		0	0	0	0	0
1	0	0		0	0	0	1	0
2	0	0		0	0	0	0	0
3	0	0		0	0	0	0	0
4	0	0		0	0	0	1	0

[5 rows x 3002 columns]

```
In [18]: # Drop the 'Email No.' column as it's just an identifier
X = df.drop(columns=['Email No.', 'Prediction']) # Drop 'Email No.' and 'Prediction' columns
y = df['Prediction'] # 'Prediction' column is the target (spam = 1, not spam = 0)
```

```
In [19]: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [20]: # Normalize the data (standardize the features for better performance with KNN)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [21]: # Initialize and train the KNN model
knn = KNeighborsClassifier(n_neighbors=5) # You can adjust 'n_neighbors' as needed
knn.fit(X_train, y_train)
```

```
Out[21]: KNeighborsClassifier
KNeighborsClassifier()
```

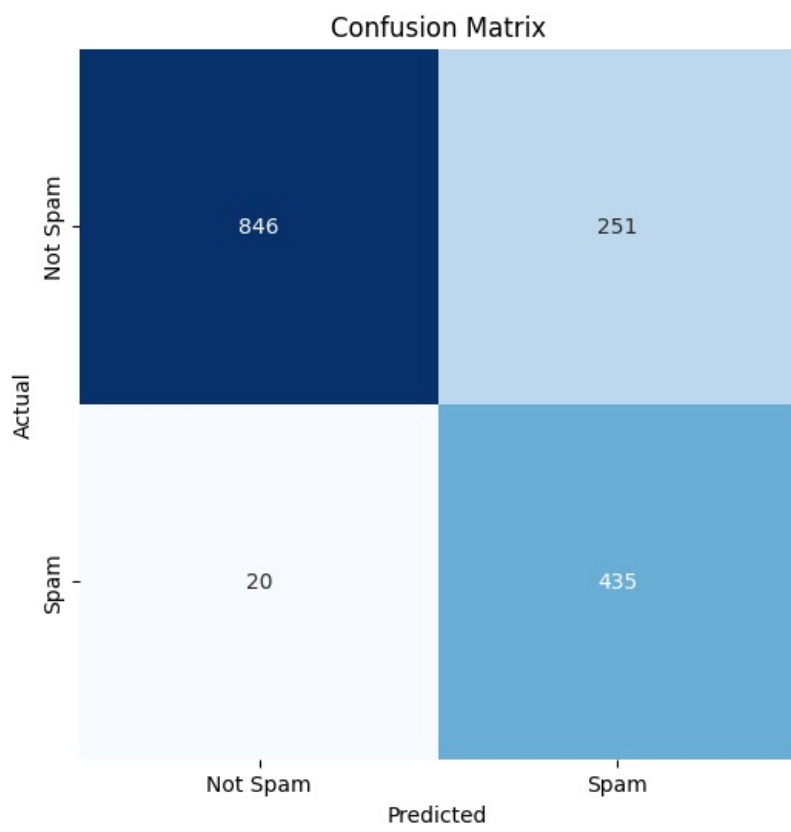
```
In [22]: # Make predictions on the test set
y_pred = knn.predict(X_test)
```

```
In [23]: # Calculate and print performance metrics
conf_matrix = confusion_matrix(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)
```

```
In [24]: # Display results
print(f"Confusion Matrix:\n{conf_matrix}")
print(f"Accuracy: {accuracy:.2f}")
```

Confusion Matrix:
[[846 251]
 [20 435]]
Accuracy: 0.83

```
In [25]: # Visualization of the confusion matrix using Seaborn
plt.figure(figsize=(6, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Not Spam', 'Spam'], yticklabels=['Not Spam', 'Spam'])
plt.title("Confusion Matrix")
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js